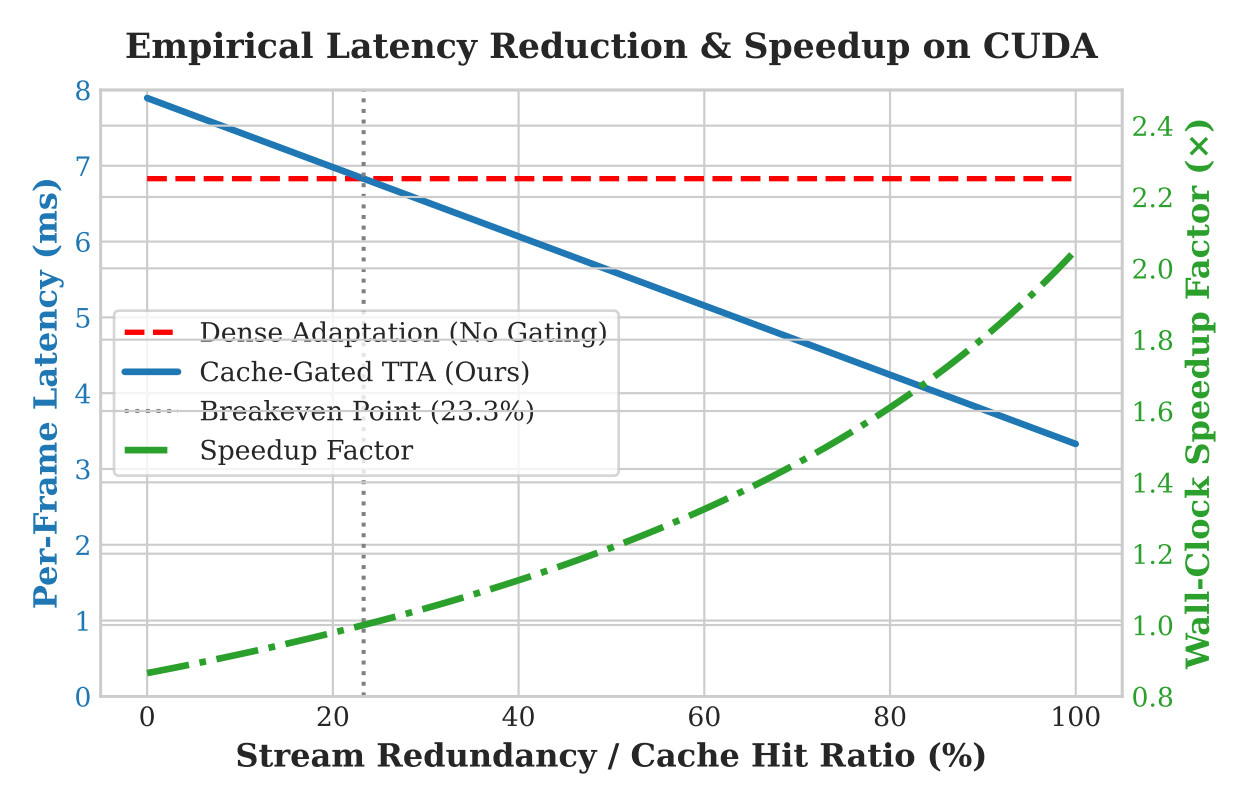
Empirical Latency Reduction & Speedup on CUDA
Per-Frame Latency (ms)
Wall-Clock Speedup Factor (×)
Stream Redundancy / Cache Hit Ratio (%)
Dense Adaptation (No Gating)
Cache-Gated TTA (Ours)
Breakeven Point (23.3%)
Speedup Factor
0
1
2
3
4
5
6
7
8
0.8
1.0
1.2
1.4
1.6
1.8
2.0
2.2
2.4
0
20
40
60
80
100